

Step 1

Refining Your Query:

- You need to get some data from the “film” table and decide to use the query `SELECT * FROM film`.
- You realize that only the “film_id” and “title” columns are needed. Write a new query that selects only those 2 columns.
- Compare the cost of the original query and the revised query, and write a few sentences explaining the comparison. Can you suggest any ways to optimize this query?

The screenshot shows the pgAdmin 4 interface. The left sidebar displays the database structure, including PostgreSQL 11, Databases (2), Rockbusters, and various objects like Casts, Catalogs, Event Triggers, Extensions (1), plpgsql, Foreign Data Wrappers, Languages, Schemas, postgres, Login/Group Roles, Tablespace, PostgreSQL 12, and PostgreSQL 13.

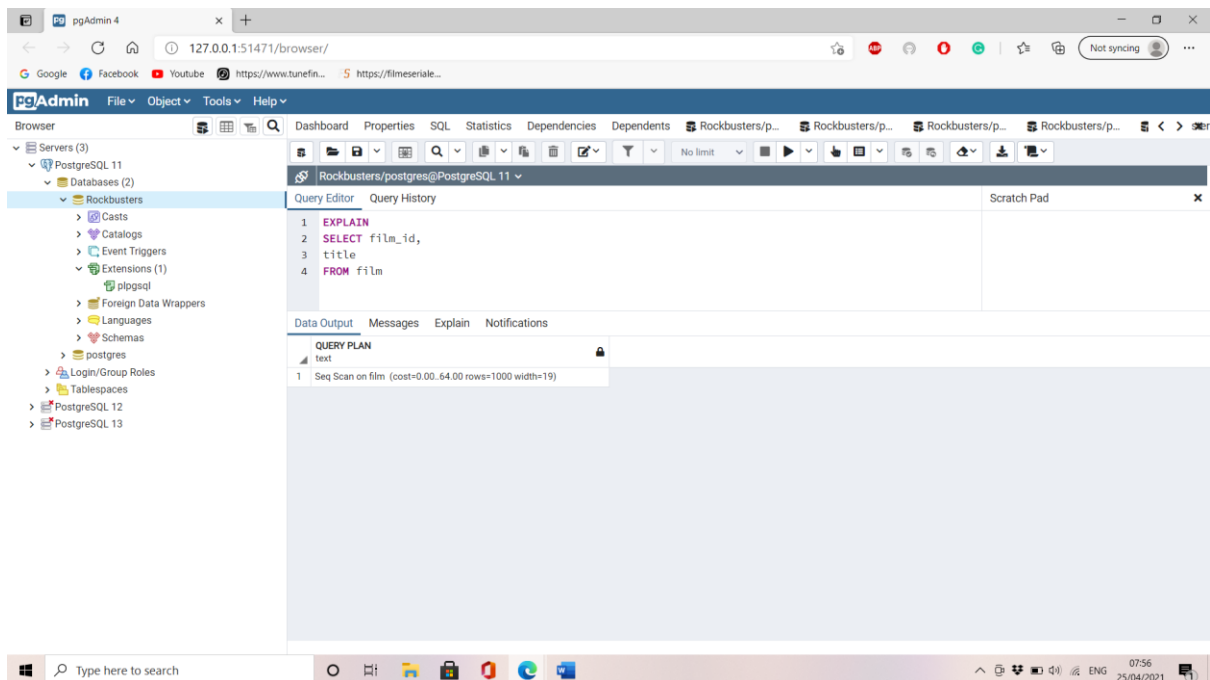
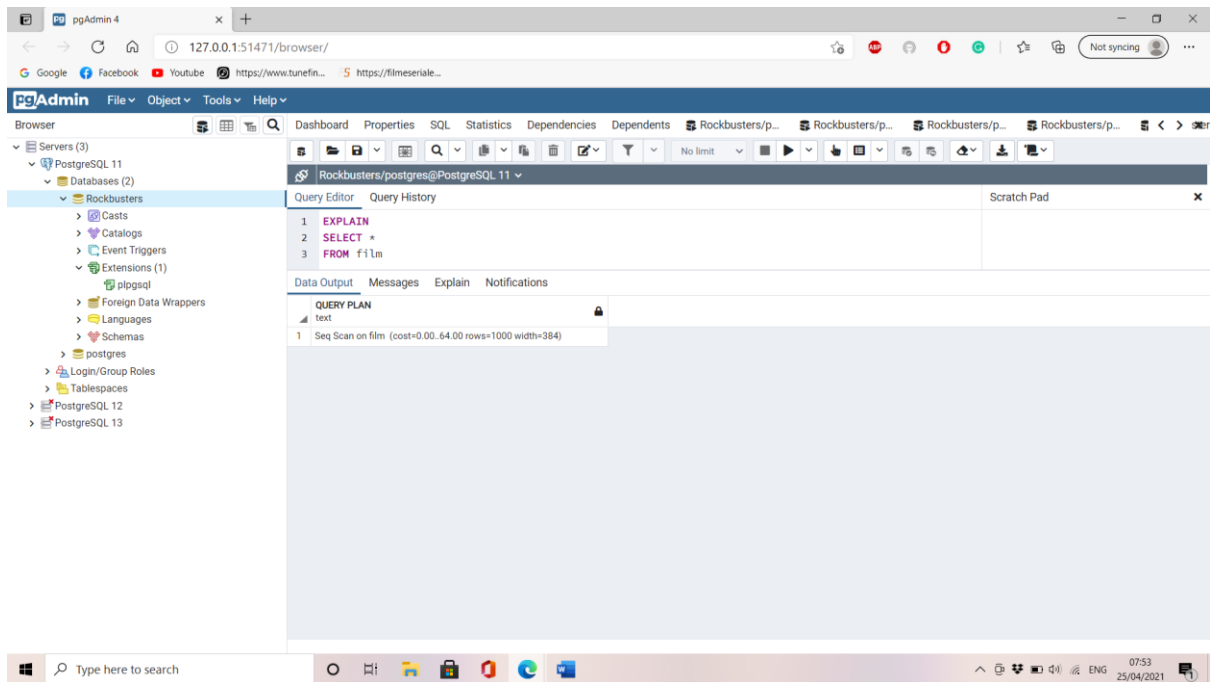
The main window shows the Query Editor with the following SQL query:

```
1 SELECT film_id,  
2 title  
3 FROM film
```

The Data Output tab displays the results of the query in a table with two columns: film_id (integer) and title (character varying (255)).

film_id	title
1	133 Chamber Italian
2	384 Grosse Wonderful
3	8 Airport Pollock
4	98 Bright Encounters
5	1 Academy Dinosaur
6	2 Ace Goldfinger
7	3 Adaptation Holes
8	4 Affair Prejudice
9	5 African Egg
10	6 Agent Truman
11	7 Airplane Sierra
12	9 Alabama Devil
13	10 Aladdin Calendar
14	11 Alamo Videotape
15	12 Alaska Phantom
16	213 Date Speed

The bottom of the screen shows the Windows taskbar with the search bar and system tray.



- The cost for the original query was 0-64 for 1000 rows and 384 columns and the cost for the revised query is 0-62 with 1000 rows and 19 columns. They are similar in cost but refining the revised query more by ordering it could save on cost a little more.

Step 2

Ordering the Data:

- In the pgAdmin Query Tool, run a query that selects every film from the “film” table, with the movies sorted by title from A to Z, then by most recent release year, and then by highest to lowest rental rate.
- Extract the data output of your query into a csv file for the film collection department to analyse in Excel. (You may need to explore how to save your output as a csv file in the Query Tool.)

The screenshot shows the pgAdmin 4 interface. The left sidebar displays the database structure, including the 'Rockbusters' database and its 'film' table. The central pane shows the 'Query Editor' with the following SQL query:

```
1 SELECT title
2 FROM film
3 ORDER BY film_id ASC
```

The 'Data Output' tab is active, displaying a table with 14 rows of film data, sorted by title. The columns are 'title' (character varying (255)) and 'release_year' (integer).

title	release_year
Academy Dinosaur	2006
Ace Goldfinger	2006
Adaptation Holes	2006
Affair Prejudice	2006
African Egg	2006
Agent Truman	2006
Airplane Sierra	2006
Airport Pollock	2006
Alabama Devil	2006
Aladdin Calendar	2006
Alamo Videotape	2006
Alaska Phantom	2006
Ali Forever	2006
Alice Fantasia	2006

The screenshot shows the pgAdmin 4 interface. The left sidebar displays the database structure, including the 'Rockbusters' database and its 'film' table. The central pane shows the 'Query Editor' with the following SQL query:

```
1 SELECT title, release_year
2 FROM film
3 ORDER BY release_year DESC
```

The 'Data Output' tab is active, displaying a table with 14 rows of film data, sorted by release year. The columns are 'title' (character varying (255)) and 'release_year' (integer).

title	release_year
Chamber Italian	2006
Grosse Wonderful	2006
Airport Pollock	2006
Bright Encounters	2006
Academy Dinosaur	2006
Ace Goldfinger	2006
Adaptation Holes	2006
Affair Prejudice	2006
African Egg	2006
Agent Truman	2006
Airplane Sierra	2006
Alabama Devil	2006
Aladdin Calendar	2006
Alamo Videotape	2006

pgAdmin 4

127.0.0.1:51471/browser/

Rockbusters/postgres@PostgreSQL 11

```

1 SELECT title, release_year, rental_rate
2 FROM film
3 ORDER BY rental_rate DESC

```

	title	release_year	rental_rate
1	French Holiday	2006	4.99
2	Bucket Brotherhood	2006	4.99
3	Frisco Forrest	2006	4.99
4	Prejudice Oleander	2006	4.99
5	Frontier Cabin	2006	4.99
6	Poseidon Forever	2006	4.99
7	Fugitive Maguire	2006	4.99
8	Wyoming Storm	2006	4.99
9	Pluto Oleander	2006	4.99
10	Platoon Instinct	2006	4.99
11	Galaxy Sweethearts	2006	4.99
12	Games Bowfinger	2006	4.99
13	Pity Bound	2006	4.99
14	Trap Guys	2006	4.99

Successfully run. Total query runtime: 336 msec. 1000 rows affected.

Step 3

Grouping Data:

- The strategy department has asked you the questions below. Write a SQL query to retrieve the correct answers, then extract your results as a csv file.
- What is the average rental rate for each rating category?
- What are the minimum and maximum rental durations for each rating category?

pgAdmin 4

127.0.0.1:51471/browser/

Rockbusters/postgres@PostgreSQL 11

```

1 SELECT rating,
2 AVG (rental_rate)
3 FROM film
4 GROUP BY rating

```

	rating	avg
1	NC-17	2.9709523809523810
2	G	2.8888764044943820
3	PG	3.0518556701030928
4	PG-13	3.0348430493273543
5	R	2.9387179487179487

pgAdmin 4

127.0.0.1:51471/browser/

pgAdmin File Object Tools Help

Browser Servers (3) PostgreSQL 11 Databases (2) Rockbusters Casts Catalogs Event Triggers Extensions (1) plpgsql Foreign Data Wrappers Languages Schemas postgres Login/Group Roles Tablespaces PostgreSQL 12 PostgreSQL 13

Query Editor Query History

```
1 SELECT rating,
2 MIN (rental_duration)
3 FROM film
4 GROUP BY rating
```

Data Output Messages Explain Notifications

rating	mpaa_rating	min	smallest
1	NC-17		3
2	G		3
3	PG		3
4	PG-13		3
5	R		3

Successfully run. Total query runtime: 369 msec. 5 rows affected.

pgAdmin 4

127.0.0.1:51471/browser/

pgAdmin File Object Tools Help

Browser Servers (3) PostgreSQL 11 Databases (2) Rockbusters Casts Catalogs Event Triggers Extensions (1) plpgsql Foreign Data Wrappers Languages Schemas postgres Login/Group Roles Tablespaces PostgreSQL 12 PostgreSQL 13

Query Editor Query History

```
1 SELECT rating,
2 MAX (rental_duration)
3 FROM film
4 GROUP BY rating
```

Data Output Messages Explain Notifications

rating	mpaa_rating	max	smallest
1	NC-17	7	
2	G	7	
3	PG	7	
4	PG-13	7	
5	R	7	

Successfully run. Total query runtime: 179 msec. 5 rows affected.

Step 4

Database Migration:

- Your team has decided to use an external tool to collect data on user behaviour in the new Rockbuster Android app. Data collected from this new source will need to be loaded into the data warehouse before you can analyse it.
- Can you outline the procedure for migrating the data and who will be responsible for it?
- The ETL procedure should be a data engineer's task but knowing the process as a data analyst would help in coordinating the process. As first step data should be extracted from different data sources (for checking user behaviour would be helpful to see, what customer search mostly, payment tracking of a specific period, most rented movies). Then it's transformed in an appropriate format and then loaded in the data warehouse.
- What problems do you foresee if you start analysing the data before it has been loaded into the data warehouse?
- Data might be incomplete or duplicate as it comes from different sources, could be involuntary manipulated.